



**The Faculty of Information and Communication Technology
Mahidol University**

**ITCS443 Parallel and Distributed Systems
Semester 1, 2023**

**Group-Project:
Parallel Merge Sort via OpenMP**

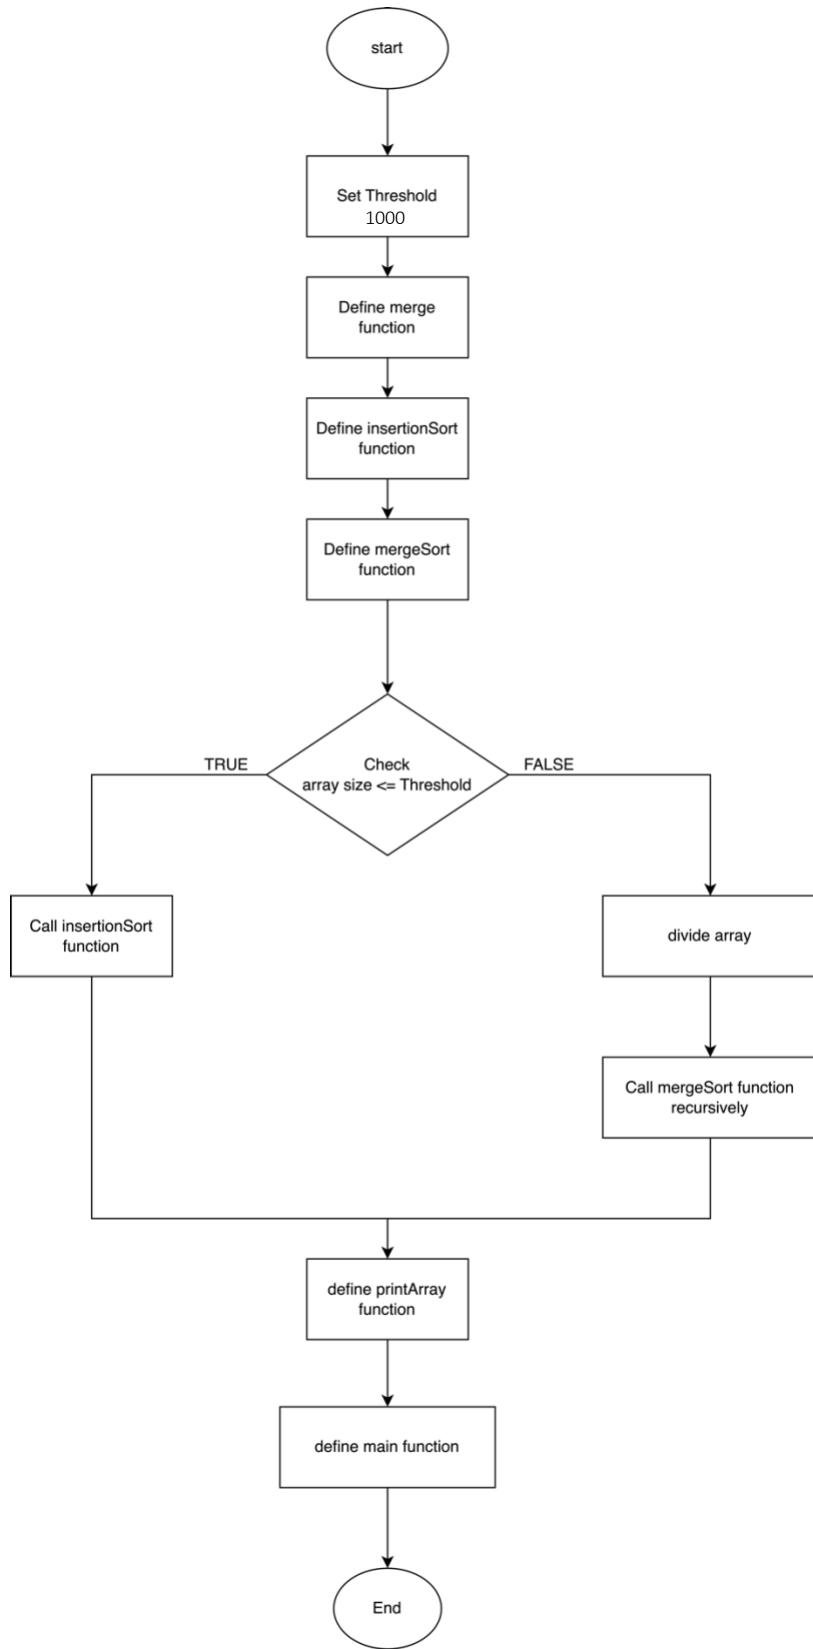
Submission Date: October 21, 2023
Section: 1
Student ID / Name: 6488179 Ponnappassorn Iamborisut
6488181 Thadeeya Duangkaew
6488210 Ravikarn Jarungjittittawas
Instructors: Dr. Sudsanguan Ngamsuriyaroj
(sudsanguan.nga@mahidol.ac.th)

Table of Contents

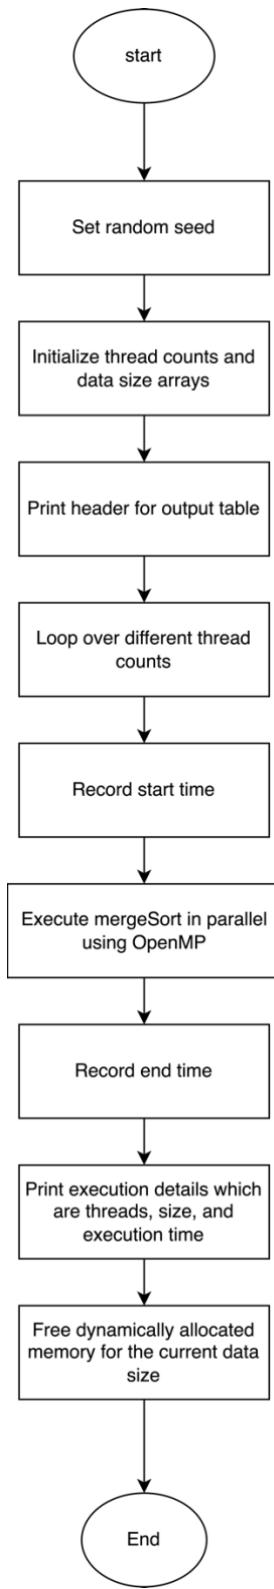
Explanation the algorithm and implementation	1
Testing result - Screenshot	7
Testing result table – Screenshot on PuTTY	17
Testing result table.....	18
Data set 10,000 random integer numbers	18
Data set 100,000 random integer numbers	18
Data set 500,000 random integer numbers	19
Data set 1,000,000 random integer numbers	19
Data set 2,000,000 random integer numbers	20
Speedup Comparison Table	20
Speed Up Graph Comparison.....	21

Explanation the algorithm and implementation

Flowchart:



In main function:



Data Structures that our group use for Merge Sort Algorithm:

- **Array:** The array to be sorted.
- **Thread count:** The number of threads to use for parallel execution.
- **Threshold:** The minimum size of a sub-array for which insertion sort will be used instead of merge sort.

Source Code:

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h> // OpenMP library for parallel programming
#include <time.h>

#define THRESHOLD 1000 // Threshold for using local sorting

// Function to merge two sub-arrays of arr[]
void merge(int arr[], int left, int mid, int right)
{
    //n1 and n2 is to calculate the sizes of two sub-arrays to be merged.
    int n1 = mid - left + 1;
    int n2 = right - mid;

    int *leftArr = (int *)malloc(n1 * sizeof(int));
    int *rightArr = (int *)malloc(n2 * sizeof(int));

    int i;
    for (i = 0; i < n1; i++)
        leftArr[i] = arr[left + i];

    int j;
    for (j = 0; j < n2; j++)
        rightArr[j] = arr[mid + 1 + j];

    //Merging the Two Sub-arrays
    i = 0;
    j = 0;
    int k;
    k = left;
    while (i < n1 && j < n2)
    {

```

```

if (leftArr[i] <= rightArr[j])
{
    arr[k] = leftArr[i];
    i++;
}
else
{
    arr[k] = rightArr[j];
    j++;
}
k++;
}
while (i < n1)

{
    arr[k] = leftArr[i];
    i++;
    k++;
}
while (j < n2)
{
    arr[k] = rightArr[j];
    j++;
    k++;
}
// Free dynamically allocated memory
free(leftArr);
free(rightArr);
}

// Function to perform insertion sort on a sub-array of arr[]
void insertionSort(int arr[], int n)
{
    int i, key, j;
    for (i = 1; i < n; i++)
    {
        // sorting by using insertion sort algorithm
        key = arr[i];
        j = i - 1;

```

```

        while (j >= 0 && arr[j] > key)
        {
            arr[j + 1] = arr[j];
            j = j - 1;
        }
        arr[j + 1] = key;
    }

}

// Function to perform merge sort on a sub-array of arr[]
void mergeSort(int arr[], int left, int right)
{
    if (left < right)
    {
        //mid index is calculated as the average of left and right, using
        integer arithmetic to prevent overflow.
        int mid = left + (right - left) / 2;

        //THRESHOLD Check for Insertion Sort:
        if (right - left + 1 <= THRESHOLD)
        {
            insertionSort(arr + left, right - left + 1);
        }
        else
        {
            #pragma omp task //executed in parallel by the OpenMP runtime.
            mergeSort(arr, left, mid);

            #pragma omp task
            mergeSort(arr, mid + 1, right);

            #pragma omp taskwait //used to wait for all the parallel tasks
            to complete before proceeding to the next step.
            merge(arr, left, mid, right);
        }
    }
}

// Function to print an array
void printArray(int arr[], int size)

```

```

{
    int i;
    printf("Sorted Data: ");
    for (i = 0; i < size; i++)
        printf("%d ", arr[i]);
    printf("\n");
}

int main()
{
    srand(123456); // Set random seed
    int num_threads[] = {1, 4, 8, 12, 16}; // Array for thread counts
    int data_sizes[] = {10000, 100000, 500000, 1000000, 2000000}; // Array
for data sizes
    printf("Data Size (N) | Num Threads | Execution Time (seconds)\n");
    printf("=====\\n");
    int i, t;
    for (i = 0; i < 5; i++) // Loop over different data sizes
    {
        int size = data_sizes[i]; // Current data size
        int *arr = (int *)malloc(size * sizeof(int)); // Array to be
sorted
        // Generate random numbers for the array
        int j;
        for (j = 0; j < size; j++)
        {
            arr[j] = rand() % 100000; // Set each element randomly to 0-
99999
        }
        for (t = 0; t < 5; t++) // Loop over different thread counts
        {
            double start_time = omp_get_wtime(); // Record start time
            #pragma omp parallel num_threads(num_threads[t]) // Parallel
execution with specified thread count
            {
                #pragma omp single // Single
directive to start parallel execution with one thread
                mergeSort(arr, 0, size - 1); // Perform merge sort in
parallel

```

```

    }

    double end_time = omp_get_wtime(); // Record end time

        // Uncomment printArray function to print data that is already
sorted by merge sort.
        //printArray(arr, size);

//printf("=====\\n");
    //printf("Data Size: %d, Thread number: %d\\n", size,
num_threads[t]);
    //printf("Execution Time (seconds): %lf\\n", end_time -
start_time);

    // Print execution details (threads, size, and execution time)

    printf("%10d      |  %6d      |      %lf\\n", size, num_threads[t],
end_time - start_time);

    if (num_threads[t] == 16)
    {

printf("=====\\n");
    }
    // Free dynamically allocated memory for this data size
    free(arr);
}

return 0;
}

```

Testing result - Screenshot

		Data set 10,000 random integer numbers			
P = 1					
P = 4					
P = 8					

$$P = 12$$

$$P = 16$$

Data set 100,000 random integer numbers

$$P = 1$$

$$P = 4$$

P = 8

$$P = 12$$

P = 16

Data set 500,000 random integer numbers

$$\underline{P} = 12$$

$$P = 16$$

Data set 1,000,000 random integer numbers

$$P = 12$$

$$P = 16$$

Data Size: 1000000, Thread number: 16
 Execution Time (seconds): 0.39185

Data set 2,000,000 random integer numbers

$$P = 12$$

P = 16

Testing result table – Screenshot on PuTTY

The screenshot shows a PuTTY terminal window with the title "10.34.73.11 - PuTTY". The terminal displays a series of numbers followed by a table of execution times for a mergesort algorithm. The table has columns for Data Size (N), Num Threads, and Execution Time (seconds). The data shows that execution time decreases as the number of threads increases for a fixed data size.

Data Size (N)	Num Threads	Execution Time (seconds)
10000	1	0.006001
10000	4	0.012585
10000	8	0.009202
10000	12	0.000332
10000	16	0.000456
100000	1	0.076635
100000	4	0.038548
100000	8	0.011089
100000	12	0.002357
100000	16	0.003035
500000	1	0.475178
500000	4	0.060750
500000	8	0.021707
500000	12	0.022296
500000	16	0.025394
1000000	1	0.958224
1000000	4	0.079994
1000000	8	0.051818
1000000	12	0.034618
1000000	16	0.032866
2000000	1	1.934395
2000000	4	0.147594
2000000	8	0.097983
2000000	12	0.106123
2000000	16	0.097725

-bash-4.1\$

Testing result table

- Data set 10,000 random integer numbers

N of Threads Trial Set \	1	4	8	12	16
Set 1	0.003106	0.007394	0.000849	0.000484	0.001022
Set 2	0.003102	0.000300	0.005816	0.000720	0.000637
Set 3	0.003169	0.000622	0.000455	0.000608	0.000671
Set 4	0.003018	0.000275	0.006441	0.000436	0.000671
Set 5	0.003013	0.000335	0.000515	0.000401	0.000447
Average	0.0030816	0.0017852	0.0028152	0.0005298	0.0006896
Speed Up	1	1.726193144	1.09462916	5.81653454	4.46867749

- Data set 100,000 random integer numbers

N of Threads Trial Set \	1	4	8	12	16
Set 1	0.042114	0.021525	0.011907	0.003952	0.002551
Set 2	0.046678	0.002902	0.002162	0.002595	0.005552
Set 3	0.054252	0.019436	0.002459	0.003482	0.002119
Set 4	0.044777	0.003131	0.004087	0.010883	0.002307
Set 5	0.059678	0.003393	0.001643	0.002580	0.002750
Average	0.0494998	0.0100774	0.0044516	0.0046984	0.0030558
Speed Up	1	4.911961419	11.1195525	10.5354589	16.1986387

Testing result table

- Data set 500,000 random integer numbers

N of Threads Trial Set \	1	4	8	12	16
Set 1	0.253251	0.009861	0.016426	0.027955	0.009160
Set 2	0.243190	0.025369	0.020490	0.10222	0.017037
Set 3	0.239614	0.009536	0.009227	0.016659	0.027199
Set 4	0.255801	0.011405	0.031367	0.010646	0.009637
Set 5	0.261955	0.010347	0.030555	0.011877	0.020454
Average	0.2507622	0.0133036	0.021613	0.0338714	0.0166974
Speed Up	1	18.84919871	11.6023782	7.40336095	15.0180387

- Data set 1,000,000 random integer numbers

N of Threads Trial Set \	1	4	8	12	16
Set 1	0.479794	0.052546	0.030082	0.034523	0.019378
Set 2	0.467432	0.060649	0.041398	0.032933	0.035166
Set 3	0.476305	0.051676	0.047858	0.036101	0.031697
Set 4	0.470280	0.067127	0.055691	0.035658	0.023004
Set 5	0.549390	0.022881	0.036992	0.026048	0.028851
Average	0.4886402	0.0509758	0.0424042	0.0330526	0.0276192
Speed Up	1	9.585728915	11.5233916	14.7837144	17.6920476

Testing result table

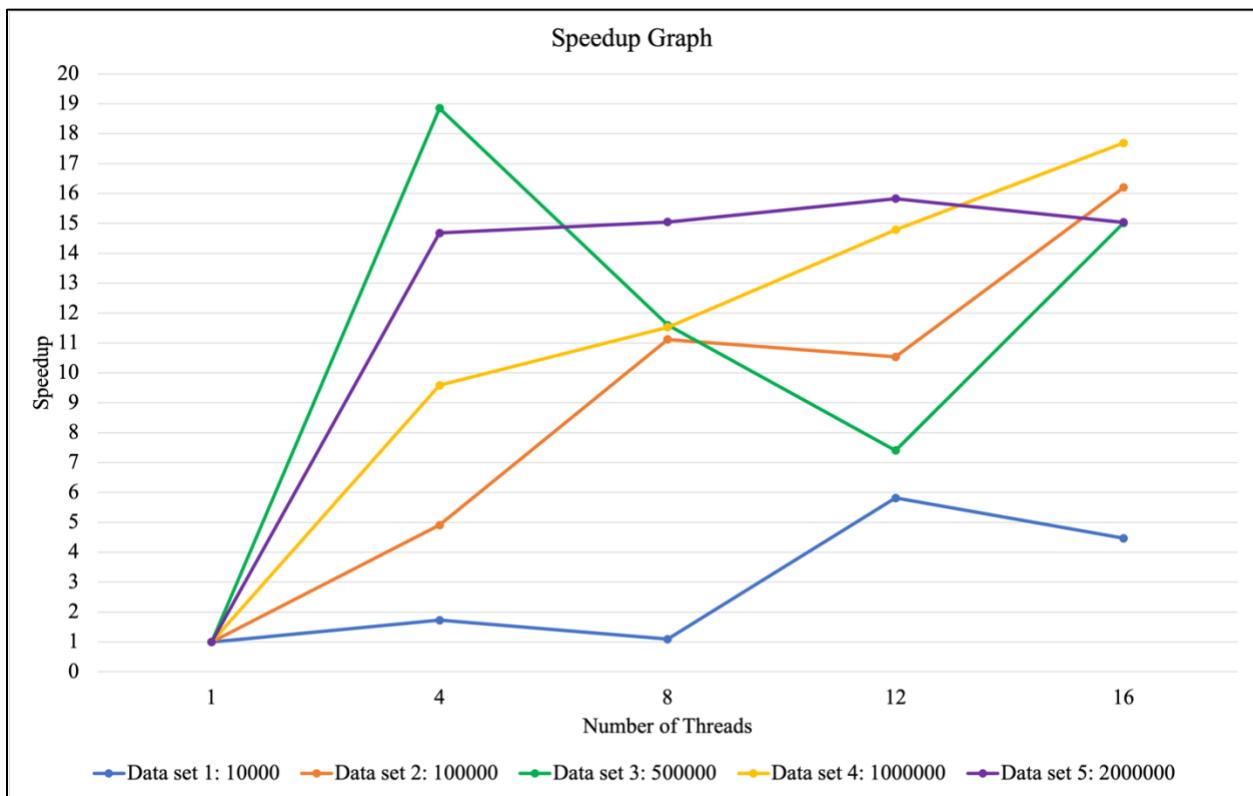
- Data set 2,000,000 random integer numbers

N of Threads Trial Set \	1	4	8	12	16
Set 1	1.022677	0.072124	0.059078	0.067891	0.072779
Set 2	1.008235	0.076407	0.057658	0.054718	0.063926
Set 3	0.986500	0.066537	0.089172	0.068558	0.088041
Set 4	1.073811	0.059825	0.053130	0.083254	0.066211
Set 5	1.053302	0.075541	0.082894	0.0505844	0.051266
Average	1.028905	0.0700868	0.0683864	0.06500108	0.0684446
Speed Up	1	14.68043911	15.0454623	15.8290447	15.0326688

Speedup Comparison Table

P	Speedup				
	1	4	8	12	16
Data set 1: 10000	1	1.726193144	1.09462916	5.81653454	4.46867749
Data set 2: 100000	1	4.911961419	11.1195525	10.5354589	16.1986387
Data set 3: 500000	1	18.84919871	11.6023782	7.40336095	15.0180387
Data set 4: 1000000	1	9.585728915	11.5233916	14.7837144	17.6920476
Data set 5: 2000000	1	14.68043911	15.0454623	15.8290447	15.0326688

Speed Up Graph Comparison



The graph shows that the speedup generally increases as the number of compute threads increases, but that the rate of speedup decreases as the number of compute threads increases. The graph also shows that the speedup varies depending on the size of the data set. For example, Data Set 5 shows the highest speedup, while Data Set 1 shows the lowest speedup. This is because larger data sets are typically more amenable to parallelization, since they can be divided into more independent tasks.